

Energy Forecast Analysis

Team 4

Ashish Dass, Anamika Jha, Shruti Narain

October 28, 2016

Abstract

In this assignment we have focused on developing a model designed for power consumption forecasting using multiple linear regressions. The scripts and the modeling have been done using R. The powers consumed for a single account based on different factors are forecasted. The forward, backward and step wise selection of variables are done to reach to an optimum model for prediction. Use of "Prune Trees" are done to get a visualization of the decision rules of dependent variables.

1. Sample Section

A part of the sample data comes from the data collected in 2014 at the Mildred school in the form of Raw-Data1.csv and Rawdata2.csv. Other half of the data has been extracted online form <https://www.wunderground.com/weather>. The data from both the sources has been combined using R script and the below modification has been done.

1.1 Data from csv files

Power: The power column contains values in 3 different units (kWh, kVARh and Power factor). The expected result was to get the power in kWh so for every hour of every single day kVARh multiplied with power factor gives kWh power consumption. Hence the value obtained from this multiplication summed with the kWh value gives us the total power consumed for that hour.

- After the power consumption has been obtained for every day, the power value given for every 5 minutes is clubbed together to obtain the value for an hour.

Now that we have a column for power consumption of every day for every hour, we have added separate columns for hour, month, day, year, day of week, weekday and peak hour all of them derived from the date and hour column using the below logic:

- Hour: has been assigned 0-23 values starting from 00:00 to 23:59 based on the values hour column we achieved after clubbing every 5 mins of data.
- Month: January to December has been assigned values 1-12 respectively derived from the date column.
- Day: values assigned from 1-31 based on the date column.
- Year: the year value derived from date column.
- Day of Week: Assigned values 0-6 from Sun-Sat respectively to values taken from date column.
- Weekday: value 1 if the day is a weekday else 0 (with use of chron library).
- Peak hour: for 7 am to 7 pm value 1 has been assigned else 0.

1.2 Data from wunderground website

The rest of the columns have been retrieved from wunderground website with the use of devtools library and "Ram-N/weatherData" github install. The data that this website provides us for Boston city is not consistent with the format that we need. Hence data cleaning has been performed to make it consistent with our already created date and time columns.

- Some erroneous value in wind_speedMPH column were replaced with 0.
- Humidity columns had N/A in the records which were replaced by 0.
- erroneous values in Wind_Direction were replaced with NA.
- The clean data was then grouped using Date and hour and mean of Humidity, Visibility, Wind Speed, Sea level pressure, Temperature, Wind Direction degrees were calculated using summarize function.
- Top most values for wind direction and Condition were considered.

The two data sets are then merged based on the hour and date columns.

2. Multiple-Linear Regression

Of all the columns we have in our SampleOutput.csv we did variable selection based on the forward, backward and step wise approach.

Backward Selection:

```
install.packages("leaps")
library(leaps)
#### Backward selection
regfit.bwd=regsubsets(power_~.,data=matWeatherFinal3,nvmax=11, method
="backward")
B=summary(regfit.bwd)
names(B)
B
B$rss
B$adjr2
coef(regfit.bwd,6)
```

```

1  ### Regression (Subset selection)
2  ### Needed package and datasets
3
4  matweatherFinal3=na.omit(matweatherFinal3) # Get rid of NAs
5  install.packages("leaps")
6  library(leaps)
7
8  #### Backward selection
9  regfit.bwd=regsubsets(power_~.,data=matweatherFinal3 ,nvmax=11, method="backward")
10
11 B=summary(regfit.bwd)
12 names(B)
13 B
14 B$rss
15 B$adjr2
16 coef(regfit.bwd,6)
17
18 par(mfrow=c(2,2))
19 plot(B$rss ,xlab="Number of variables ",ylab="RSS", type="l")
20 plot(B$adjr2 ,xlab="Number of Variables ", ylab="Adjusted Rsq",type="l")
21 coef(regfit.bwd ,6)
22

```

18:1 (Top Level) R Script

```

Console E:/Sem4-Fall 2016/ADS/Assignment 2 Materials/
> ( 1 ) " " " "
10 ( 1 ) " " " "
11 ( 1 ) " " " "
12 ( 1 ) " " " "
> B$rss
[1] 462880793 443825151 423655917 402579078 379305392 356678985 336713251 316293857
[9] 297164183 278849873 265368671 254489824
> B$adjr2
[1] 0.03777876 0.07728140 0.11910887 0.16283365 0.21113773 0.25810695 0.29955247
[8] 0.34195168 0.38167740 0.41971580 0.44770440 0.47028288
> coef(regfit.bwd,6)
(Intercept)      hour__8      hour__9      hour__10      hour__11      hour__12      hour__13
    364.8176      262.6524      278.5281      287.2542      291.9627      282.7587      275.0793

```

- The first mean squared value comes from all the variables selected.
- From step 2, one variable is being removed at a time.
- The variables are dropped one by one till we reach our best result.

Sno	Mean R squared	Removed	Equation
Step1	1	0.7211	power_~ hour__ + month__ + day__ + dayofWeek__ + Weekday__ + Wind_Direction + Conditions + Temperature + Dew_PointF + Humidity + Sea_Level_PressureIn + VisibilityMPH + Wind_SpeedMPH + WindDirDegrees + peakOfhour__
Step 2	1	0.6308	power_~ hour__ + month__ + day__ + dayofWeek__ + Wind_Direction + Conditions + Temperature + Dew_PointF + Humidity + Sea_Level_PressureIn + VisibilityMPH + Wind_SpeedMPH + WindDirDegrees + peakOfhour__
	2	0.6797	month__
	3	0.7145	day__
	4	0.7069	dayofWeek__
	5	0.7211	Weekday__
	6	0.7171	Wind_Direction
	7	0.7189	Conditions
	8	0.721	Temperature
	9	0.7205	Dew_PointF
	10	0.7198	Humidity
	11	0.7166	Sea_Level_PressureIn
	12	0.7211	VisibilityMPH
	13	0.721	Wind_SpeedMPH
	14	0.7211	WindDirDegrees
	15	0.7211	peakOfhour__

Step 3	1	0.6308	hour__
	2	0.6797	month__
	3	0.7145	day__
	4	0.6046	dayofWeek__
	5	0.7171	Wind_Direction
	6	0.7189	Conditions
	7	0.721	Temperature
	8	0.7205	Dew_PointF
	9	0.7198	Humidity
	10	0.7211	Sea_Level_PressureIn
	11	0.721	VisibilityMPH
	12	0.721	Wind_SpeedMPH
	13	0.7211	WindDirDegrees
	14	0.7211	peakOfhour__
Step 4	1	0.6308	hour__
	2	0.679	month__
	3	0.7145	day__
	4	0.6046	dayofWeek__
	5	0.7171	Wind_Direction
	6	0.7189	Conditions
	7	0.7209	Temperature
	8	0.7204	Dew_PointF

Summary of the most optimum model obtained using backward selection on train data:

```

1 install.packages("car")
2 library(car)
3 samp <- floor(0.75* nrow(matweatherFinal))
4 samp
5 set.seed(123)
6 training.index <- sample(seq_len(nrow(matweatherFinal)), size= samp)
7 train <- matweatherFinal[training.index, ]
8 test <- matweatherFinal[-training.index, ]
9 h<-matweatherFinal
10
11 na.omit(test)
12 na.omit(train)
13 #check sea level and temperature collinearity
14 lm.fit = lm(power_ ~ hour__ + month__ + day__ + dayofweek__ +
15             wind_direction + Temperature + Dew_PointF +
16             Humidity + Sea_Level_PressureIn ,data=train)
17 summary(lm.fit)
18 RegressionOutput <- summary(lm.fit)$coefficients[,1]
19 write.csv(RegressionOutput, file = "RegressionOutput1.csv")
20
21 vif(lm.fit)
22 install.packages("forecast")

```

24:26 (Top Level) R Script

```

Console: E:/Sem4-Fall 2016/ADS/Assignment 2 Materials/
wind_direction 21.2707 11.2307 2.043 0.004425
Temperature -1.8866 1.0855 -1.738 0.082239 .
Dew_PointF 4.3813 1.1648 3.761 0.000170 ***
Humidity -2.9834 0.5135 -5.810 6.53e-09 ***
Sea_Level_PressureIn -15.5664 8.3306 -1.869 0.061727 .
---
Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

Residual standard error: 128.2 on 6426 degrees of freedom
(52 observations deleted due to missingness)
Multiple R-squared:  0.7177,    Adjusted R-squared:  0.7137
F-statistic: 179.5 on 91 and 6426 DF,  p-value: < 2.2e-16

```

Regression output:

```
4 samp
5 set.seed(123)
6 training.index <- sample(seq_len(nrow(matweatherFinal)), size= samp)
7 train <- matweatherFinal[training.index, ]
8 test <- matweatherFinal[-training.index, ]
9 h<-matweatherFinal
10
11 na.omit(test)
12 na.omit(train)
13 #check sea level and temperature collinearity
14 lm.fit = lm(power_ ~ hour_ + month_ + day_ +dayofweek_ +
15             wind_Direction + Temperature + Dew_PointF +
16             Humidity + Sea_Level_PressureIn ,data=train)
17 summary(lm.fit)
18 RegressionOutput <- summary(lm.fit)$coefficients[,1]
19 write.csv(RegressionOutput, file = "RegressionOutput1.csv")
20
21 vif(lm.fit)
22 install.packages("forecast")
23 library(forecast)
24 pred=predict(lm.fit,test)
25 PerformanceMetrics <- accuracy(pred,train$power_)
```

18:17 (Top Level) R Script

Console E:/Sem4-Fall 2016/ADS/Assignment 2 Materials/

Multiple R-squared: 0.7177, Adjusted R-squared: 0.7137
F-statistic: 179.5 on 91 and 6426 DF, p-value: < 2.2e-16

```
> RegressionOutput <- summary(lm.fit)$coefficients[,1]
> RegressionOutput
      (Intercept)      hour__1      hour__2
669.1230718      2.6196710     13.4885581
      hour__3      hour__4      hour__5
10.0119812     17.9064844     37.0014994
      hour__6      hour__7      hour__8
236.7689732     356.0020129     378.0166407
      hour__9      hour__10     hour__11
390.5932896     392.3482593     394.3028698
      hour__12     hour__13     hour__14
382.4697477     369.8284938     334.8893453
```

Performance matrix:

```
11 na.omit(test)
12 na.omit(train)
13 #check sea level and temperature collinearity
14 lm.fit = lm(power_ ~ hour_ + month_ + day_ +dayofweek_ +
15             wind_Direction + Temperature + Dew_PointF +
16             Humidity + Sea_Level_PressureIn ,data=train)
17 summary(lm.fit)
18 RegressionOutput <- summary(lm.fit)$coefficients[,1]
19 write.csv(RegressionOutput, file = "RegressionOutput1.csv")
20
21 vif(lm.fit)
22 install.packages("forecast")
23 library(forecast)
24 pred=predict(lm.fit,test)
25 PerformanceMetrics <- accuracy(pred,train$power_)
26
27 write.csv(t(PerformanceMetrics), file = "PerformanceMetrics1.csv")
28
29 lm.fit1 = lm(power_ ~ ., data=train)
30 summary(lm.fit1)
31
32 pred=predict(lm.fit1,test)
```

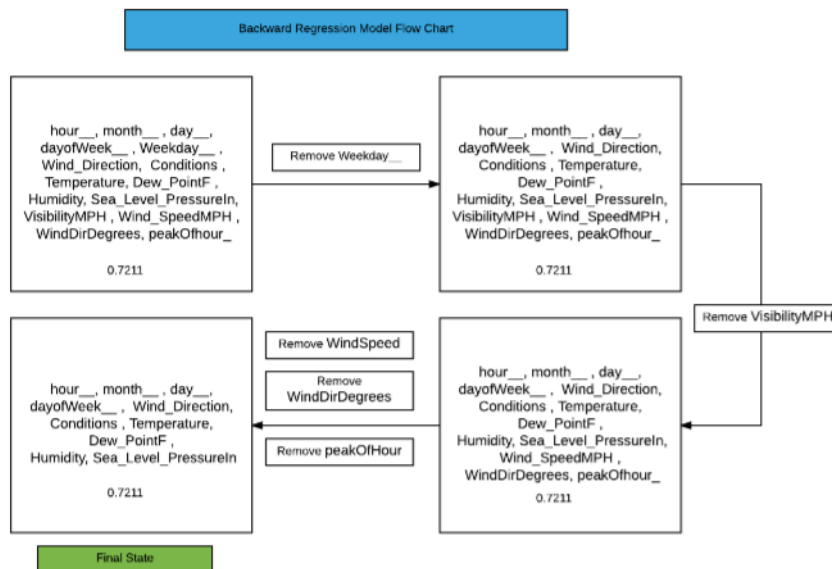
25:9 (Top Level) R Script

Console E:/Sem4-Fall 2016/ADS/Assignment 2 Materials/

```
wind_DirectionSE      wind_DirectionSE      wind_DirectionSE      wind_DirectionSE
39.8053023      5.1920003      17.2130871
wind_DirectionSE      Temperature      Dew_PointF
31.9707317      -1.8866461      4.3812549
Humidity      Sea_Level_PressureIn
-2.9833748      -15.5663799
```

```
> t(PerformanceMetrics)
Test set
ME      -14.29269
RMSE     982.01770
MAE     267.09981
MPE     -31.12927
MAPE     71.48662
>
```

Flowchart:



Step wise Selection:

Summary of the most optimum model obtained using step wise selection on train data:

```

1 install.packages("car")
2 library(car)
3 samp <- floor(0.75 * nrow(matweatherFinal))
4 samp
5 set.seed(123)
6 training.index <- sample(seq_len(nrow(matweatherFinal)), size= samp)
7 train <- matweatherFinal[training.index, ]
8 test <- matweatherFinal[-training.index, ]
9 h<-matweatherFinal
10
11 na.omit(test)
12 na.omit(train)
13 #check sea level and temperature collinearity
14 lm.fit = lm(power_ ~ Temperature + month___ + day___ + hour___ + dayofweek___, data=train)
15 summary(lm.fit)
16 RegressionOutput <- summary(lm.fit)$coefficients[,1]
17 write.csv(RegressionOutput, file = "RegressionOutputForward.csv")
18
19 vif(lm.fit)
20 install.packages("forecast")
21 library(forecast)
22 <
  
```

16:1 (Top Level) R Script

Console E:/Sem4-Fall 2016/ADS/Assignment 2 Materials/

```

dayofweek___3 232.3008    6.0787  38.215 < 2e-16 ***
dayofweek___4 226.2204    6.0986  37.094 < 2e-16 ***
dayofweek___5 207.5756    6.1180  33.929 < 2e-16 ***
dayofweek___6 104.9973    6.1347  17.115 < 2e-16 ***
---
Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

Residual standard error: 129.7 on 6446 degrees of freedom
(52 observations deleted due to missingness)
Multiple R-squared:  0.7102,    Adjusted R-squared:  0.707
F-statistic: 222.5 on 71 and 6446 DF,  p-value: < 2.2e-16
  
```

Performance Matrix:

```
8 test <- matweatherFinal[-training.index, ]
9 h<-matweatherFinal
10
11 na.omit(test)
12 na.omit(train)
13 #check sea level and temperature collinearity
14 lm.fit = lm(power_ ~ Temperature + month_ + day_ + hour_ +dayofweek_, data=train)
15 summary(lm.fit)
16 RegressionOutput <- summary(lm.fit)$coefficients[,1]
17 write.csv(RegressionOutput, file = "RegressionOutputStepwise.csv")
18
19 vif(lm.fit)
20 install.packages("forecast")
21 library(forecast)
22 pred=predict(lm.fit,test)
23 PerformanceMetrics <- accuracy(pred,train$power_)
24
25 write.csv(t(PerformanceMetrics), file = "PerformanceMetricsForward.csv")
26
27 lm.fit1 = lm(power_ ~ ., data=train)
28 summary(lm.fit1)
29 <
```

21:2 (Top Level) R Script

Console E:/Sem4-Fall 2016/ADS/Assignment 2 Materials/

```
month_ 3.732408 11 1.002777
day_ 1.324796 30 1.004699
hour_ 1.243753 23 1.004753
dayofweek_ 1.246360 6 1.018522
> pred=predict(lm.fit,test)
> PerformanceMetrics <- accuracy(pred,train$power_)
> t(PerformanceMetrics)
      Test set
ME    9.321017
RMSE 343.005441
MAE  249.902608
MPE  -27.334111
MAPE  68.733958
```

Regression Output:

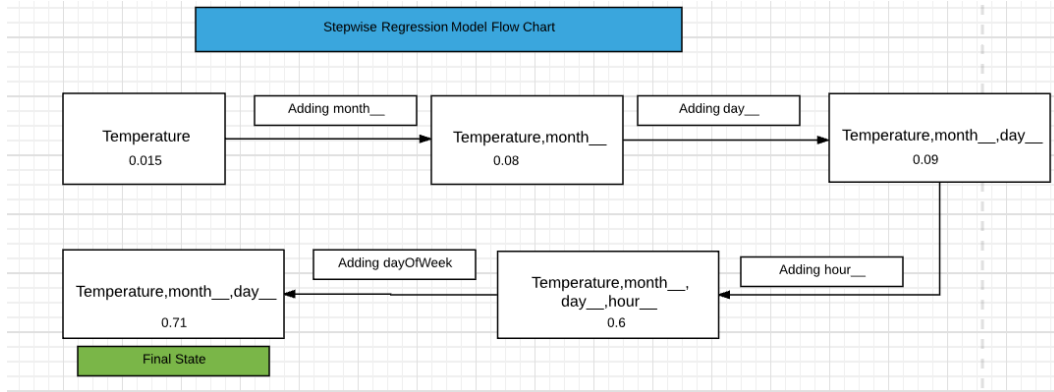
```
4 samp
5 set.seed(123)
6 training.index <- sample(seq_len(nrow(matweatherFinal)), size= samp)
7 train <- matweatherFinal[training.index, ]
8 test <- matweatherFinal[-training.index, ]
9 h<-matweatherFinal
10
11 na.omit(test)
12 na.omit(train)
13 #check sea level and temperature collinearity
14 lm.fit = lm(power_ ~ Temperature + month_ + day_ + hour_ +dayofweek_, data=train)
15 summary(lm.fit)
16 RegressionOutput <- summary(lm.fit)$coefficients[,1]
17 write.csv(RegressionOutput, file = "RegressionOutputStepwise.csv")
18
19 vif(lm.fit)
20 install.packages("forecast")
21 library(forecast)
22 pred=predict(lm.fit,test)
23 PerformanceMetrics <- accuracy(pred,train$power_)
24
25 <
```

17:61 (Top Level) R Script

Console E:/Sem4-Fall 2016/ADS/Assignment 2 Materials/

```
Residual standard error: 129.7 on 6446 degrees of freedom
(52 observations deleted due to missingness)
Multiple R-squared:  0.7102,    Adjusted R-squared:  0.707
F-statistic: 222.5 on 71 and 6446 DF,  p-value: < 2.2e-16
> RegressionOutput <- summary(lm.fit)$coefficients[,1]
> RegressionOutput
(Intercept) Temperature month_2 month_3 month_4 month_5
-24.4577501  1.7877639 -25.1623248 -22.0503418 -77.0812015 -89.8381607
month_6 month_7 month_8 month_9 month_10 month_11
-18.7332802  50.6586782  86.0343823 -46.9085329 -97.1714806 -63.0722209
month_12 day_2 day_3 day_4 day_5 day_6
-81.0270548  79.5594709  39.2386569  15.5880244  54.3827346  36.2389357
day_7 day_8 day_9 day_10 day_11 day_12
```


Flowchart:



Forward Selection:

```

install.packages("ISLR")
require(leaps)
install.packages("glmnet")
library(glmnet)
library(ISLR)
regfit.fwd = regsubsets(power_~., data= matWeatherFinal, nvmax=11, method=
"forward")
F= summary(regfit.fwd)
names(F)

```

F

```

72 #####
73
74 install.packages("ISLR")
75 require(leaps)
76 install.packages("glmnet")
77 library(glmnet)
78 library(ISLR)
79 regfit.fwd = regsubsets(power_~., data= matWeatherFinal, nvmax=11, method= "forward")
80 F= summary(regfit.fwd)
81 names(F)
82 F
83 F$rss
84
85
84:1 temp + wind speed +
R Script

```

```

> F= summary(regfit.fwd)
> names(F)
[1] "which" "rsq" "rss" "adjr2" "cp" "bic" "outmat" "obj"
> F$rsq
[1] 284797282 235593839 231020456 215512463 211910833 209211759 207751028 206798503 205994711
[10] 205172281 204403492 203693902
> F$adjr2
[1] 0.4079729 0.5101972 0.5196482 0.5518401 0.5592774 0.5648391 0.5678261 0.5697564 0.5713777
[10] 0.5730382 0.5745875 0.5760139
> coef(regfit.fwd,12)
(Intercept)          hour__      matweatherDate_
2.570388e+05      1.321335e+00      -1.600208e+01
      month__          day__      weekday__
4.879432e+02      1.437390e+01      1.720657e+02
wind_directionWSW  wind_Directionsw  ConditionsScattered Clouds
-2.919478e+00      4.589843e+00      9.982477e+01
ConditionsMostly Cloudy      Humidity      wind_SpeedMPH
6.323488e+01      -1.562453e+00      -2.546379e-04
ConditionsMist
0.000000e+00

```

- Started with Temperature alone as a dependent variable.
- Repeated the process for all the dependent variables.
- We kept adding one more variable at a time to check any improvement in the summary results.
- These steps were repeated till we reached our best results.

Example 1:

```
fit <- lm(power_ ~ Temperature , data=matweatherFinal)

summary(fit)

Call:
lm(formula = power_ ~ Temperature, data = matweatherFinal)

Residuals:
    Min       1Q   Median       3Q      Max
-312.4 -185.9 -113.6  180.1 2423.9

Coefficients:
            Estimate Std. Error t value Pr(>|t|)
(Intercept) 394.66679    4.11112   96.0   <2e-16 ***
Temperature   0.80414    0.06331   12.7   <2e-16 ***
---
Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

Residual standard error: 236.8 on 8419 degrees of freedom
Multiple R-squared:  0.0188,    Adjusted R-squared:  0.01869
F-statistic: 161.3 on 1 and 8419 DF,  p-value: < 2.2e-16
```

Example 2:

```
fit <- lm(power_ ~ Temperature , data=matweatherFinal)

summary(fit)
```

```

Call:
lm(formula = power_ ~ Dew_PointF, data = matweatherFinal)

Residuals:
    Min       1Q   Median       3Q      Max
-309.7 -186.3 -115.3  180.4 1587.8

Coefficients:
            Estimate Std. Error t value Pr(>|t|)
(Intercept)  414.26357    3.51105  117.989  <2e-16 ***
Dew_PointF    0.55620    0.06254   8.894  <2e-16 ***
---
Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

Residual standard error: 237.9 on 8419 degrees of freedom
Multiple R-squared:  0.009308, Adjusted R-squared:  0.009191
F-statistic: 79.1 on 1 and 8419 DF, p-value: < 2.2e-16

```

Summary of the most optimum model obtained using forward selection on train data:

```

1  install.packages("car")
2  library(car)
3  samp <- floor(0.75* nrow(matweatherFinal))
4  samp
5  set.seed(123)
6  training.index <- sample(seq_len(nrow(matweatherFinal)), size= samp)
7  train <- matweatherFinal[training.index, ]
8  test <- matweatherFinal[-training.index, ]
9  h<-matweatherFinal
10
11 na.omit(test)
12 na.omit(train)
13 #check sea level and temperature collinearity
14 lm.fit = lm.fit = lm(power_ ~ Temperature + month_ + day_ + hour_ +dayofweek_
15 summary(lm.fit)
16 RegressionOutput <- summary(lm.fit)$coefficients[,1]
17 write.csv(RegressionOutput, file = "RegressionOutputForward.csv")
18
19 vif(lm.fit)
20 install.packages("forecast")
21 library(forecast)
22 <

```

16:1 (Top Level) R Script

```

Console E:/Sem4-Fall 2016/ADS/Assignment 2 Materials/
dayofweek_3 226.6737 6.0962 37.183 < 2e-16 ***
dayofweek_4 226.6737 6.0962 37.183 < 2e-16 ***
dayofweek_5 207.2720 6.1146 33.898 < 2e-16 ***
dayofweek_6 104.5757 6.1320 17.054 < 2e-16 ***
Sea_Level_PressureIn 23.3449 7.4471 3.135 0.001728 **
---
Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

Residual standard error: 129.6 on 6445 degrees of freedom
(52 observations deleted due to missingness)
Multiple R-squared:  0.7106, Adjusted R-squared:  0.7074
F-statistic: 219.8 on 72 and 6445 DF, p-value: < 2.2e-16
> |

```

Regression Output:

```
6 training.index <- sample(seq_len(nrow(matweatherFinal)), size= samp)
7 train <- matweatherFinal[training.index, ]
8 test <- matweatherFinal[-training.index, ]
9 h<-matweatherFinal
10
11 na.omit(test)
12 na.omit(train)
13 #check sea level and temperature collinearity
14 lm.fit = lm.fit = lm(power_ ~ Temperature + month__ + day__ + hour__ +dayofweek__
15 summary(lm.fit)
16 RegressionOutput <- summary(lm.fit)$coefficients[,1]
17 write.csv(RegressionOutput, file = "RegressionOutputForward.csv")
18
19 vif(lm.fit)
20 install.packages("forecast")
21 library(forecast)
22 pred=predict(lm.fit,test)
23 PerformanceMetrics <- accuracy(pred,train$power_)
24
25 write.csv(t(PerformanceMetrics), file = "PerformanceMetricsForward.csv")
26
27
28
```

16:17 (Top Level) ↕ R Script

Console E:/Sem4-Fall 2016/ADS/Assignment 2 Materials/ ↕

```
> RegressionOutput <- summary(lm.fit)$coefficients[,1]
> RegressionOutput
      (Intercept)      Temperature      month__2      month__3
      -733.7741813      2.0649232      -26.2435660      -23.8853171
      month__4      month__5      month__6      month__7
      -84.0064189      -98.7122738      -30.1017513      38.7935028
      month__8      month__9      month__10      month__11
      73.5675632      -60.2127855      -103.5471478      -67.1640057
      month__12      day__2      day__3      day__4
      -86.4487545      81.3099195      41.6924738      17.2645892
      day__5      day__6      day__7      day__8
      56.6081316      38.2805195      47.5385017      50.8267705
      day__9      day__10      day__11      day__12
```

Performance Matrix:

```

10
11 na.omit(test)
12 na.omit(train)
13 #check sea level and temperature collinearity
14 lm.fit = lm.fit = lm(power_ ~ Temperature + month_ + day_ + hour_
15 summary(lm.fit)
16 RegressionOutput <- summary(lm.fit)$coefficients[,1]
17 write.csv(RegressionOutput, file = "RegressionOutputForward.csv")
18
19 vif(lm.fit)
20 install.packages("forecast")
21 library(forecast)
22 pred=predict(lm.fit,test)
23 PerformanceMetrics <- accuracy(pred,train$power_)
24
25 write.csv(t(PerformanceMetrics), file = "PerformanceMetricsForward.csv")
26
27 lm.fit1 = lm(power_ ~ ., data=train)
28 summary(lm.fit1)
29
30 pred=predict(lm.fit1,test)
31 <

```

27:1 (Top Level) ↕

```

Console E:/Sem4-Fall 2016/ADS/Assignment 2 Materials/ ↗
warning in install.packages :
package 'forecast' is in use and will not be installed
> library(forecast)
> pred=predict(lm.fit,test)
> PerformanceMetrics <- accuracy(pred,train$power_)
> t(PerformanceMetrics)
      Test set
ME      45.69453
RMSE 1860.27492
MAE    286.47067
MPE   -21.71344
MAPE    74.42619
> write.csv(t(PerformanceMetrics), file = "PerformanceMetricsForward.csv")
~ |

```

3. Prediction based on forecast data

- The Sample output is generated from Raw datasets merged, after data cleaning and feature selection (after trying Forward, Backward and Stepwise) techniques.
- The feature selection, led us to split the Sample Output into 75% training data set and 25% testing dataset, after trying a possible combinations of 60-40, 80-20, etc.
- Then after fitting the data into the multiple linear regression, we used package "Forecast" to use its predict() function to predict the data on the clean Forecast data missing on the Power/kwh(which was to be predicted).

The summary and the performance matrix of the final model selected using step wise regression

```

fDataNew * Run as 3rd - COMPLETE_PART3_A2_CL... * forecastPrediction * Run as 4th - COMPLETE_FORECASTIN... * *
Source on Save Run Source
21
22 ##### CHECKING FIT OF 1 - THE BEST MODEL #####
23
24 lm.fit = lm(kwh ~ Temperature + month + day + `Day of week` + hour,data=train) # removed pe
25 summary(lm.fit)
26
27 ##### Getting all coefficients #####
28 summary(lm.fit)$coefficients[,]
29
30
29:1 Getting all coefficients R Script
Console C:/Users/310250154/Desktop/ADS/Assignments/assignment 2/Assignment 2(1)/
4482 scattered clouds 240.00000
7251 Mostly cloudy 310.00000
4532 Mostly cloudy 110.00000
7420 Light Rain 100.00000
4811 Mostly cloudy 100.00000
3467 Mostly cloudy 220.00000
7774 Partly cloudy 360.00000
5804 Overcast 240.00000
3383 Overcast 135.00000
151 Overcast 350.00000
[ reached getoption("max.print") -- omitted 5771 rows ]
> lm.fit = lm(kwh ~ Temperature + month + day + `Day of week` + hour,data=train) # removed peakh
our and weekday #humidity is making a diff of 0.001
> summary(lm.fit)

Call:
lm(formula = kwh ~ Temperature + month + day + `Day of week` +
    hour, data = train)

Residuals:
    Min       1Q   Median       3Q      Max
-572.04  -67.31   -4.41   57.08  737.76

Coefficients:
            Estimate Std. Error t value Pr(>|t|)
(Intercept)  33.80554   13.59021    2.487  0.012890 *
Temperature    0.06962    0.03822    1.822  0.068548 .
month2     -22.52104    7.92725   -2.841  0.004512 **
month3     -11.71157    7.74514   -1.512  0.130553
month4     -48.37865    7.92839   -6.102  1.11e-09 ***
month5     -36.56030    7.73785   -4.725  2.35e-06 ***

hour20      126.29005    11.21701   11.259 < 2e-16 ***
hour21       34.66418    11.08578    3.127  0.001774 **
hour22       12.87687    11.14207    1.156  0.247847
hour23       11.40202    11.17354    1.020  0.307554

Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

Residual standard error: 129.2 on 6446 degrees of freedom
(52 observations deleted due to missingness)
Multiple R-squared:  0.7141,    Adjusted R-squared:  0.711
F-statistic: 226.8 on 71 and 6446 DF,  p-value: < 2.2e-16

>

```

- with the use of `vif()` [variance inflation factors] function we have determined the collinearity between dependent variables:

- Values more than 5 denotes collinearity between the variables. In our case all values are between 0-1 hence the dependent variables we have selected are no collinear.

```

28 summary(lm.fit)$coefficients[,1]
29
30 ##### Getting 1 coefficient #####
31 coef(summary(lm.fit))["month__02","Estimate"]
32 library("car")
33 vif(lm.fit)
34
35 pred=predict(lm.fit,test)
36 library("forecast")
37 accuracy(pred,train$kwH)
38
39 ##### CHECKING FIT OF ALL THE VARIABLES #####
40
41 lm.fit1 = lm(power_ ~ ., data=train)
42 summary(lm.fit1)
43
44 pred=predict(lm.fit1,test)
45 accuracy(pred,train$kwH)
46
47 ##### PREDICTING THE FORECAST#####
48
49

```

35:1 # Getting 1 coefficient

Console C:/Users/310250154/Desktop/ADS/Assignments/assignment 2/Assignment 2(1)/

signif. codes: 0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

Residual standard error: 129.2 on 6446 degrees of freedom

(52 observations deleted due to missingness)

Multiple R-squared: 0.7141, Adjusted R-squared: 0.711

F-statistic: 226.8 on 71 and 6446 DF, p-value: < 2.2e-16

```

> vif(lm.fit)

```

	GVIF	Df	GVIF^(1/(2*Df))
Temperature	1.167016	1	1.080285
Month	1.210154	11	1.008708
Day	1.267092	30	1.003953
Day of week	1.216142	6	1.016441
Hour	1.052748	23	1.001118

```

25 summary(lm.fit)
26
27 ##### Getting all coefficients #####
28 summary(lm.fit)$coefficients[,]
29
30 ##### Getting 1 coefficient #####
31 coef(summary(lm.fit))["month__02", "Estimate"]
32 library("car")
33 vif(lm.fit)
34
35 pred=predict(lm.fit,test)
36 library("forecast")
37 accuracy(pred,train$kw)
38
39 ##### CHECKING FIT OF ALL THE VARIABLES #####
40
41 lm.fit1 = lm(power_ ~ ., data=train)
42 summary(lm.fit1)
43
44 pred=predict(lm.fit1,test)
45 accuracy(pred,train$kw)
46
47 ##### PREDICTING THE FORECAST#####
48
49
50 pred=predict(lm.fit,forecastPrediction)
51
39:1 # CHECKING FIT OF ALL THE VARIABLES

```

Console C:/Users/310250154/Desktop/ADS/Assignments/assignment 2/Assignment 2(1)/

```

accuracy(pred,train$kw)

```

	ME	RMSE	MAE	MPE	MAPE
test set	0.5016031	313.4098	246.9187	-29.33255	68.78474

Prune Tree:

Prune Tree is used to avoid overfitting the data. The size of the tree can be selected in order to minimize the cross validated error.

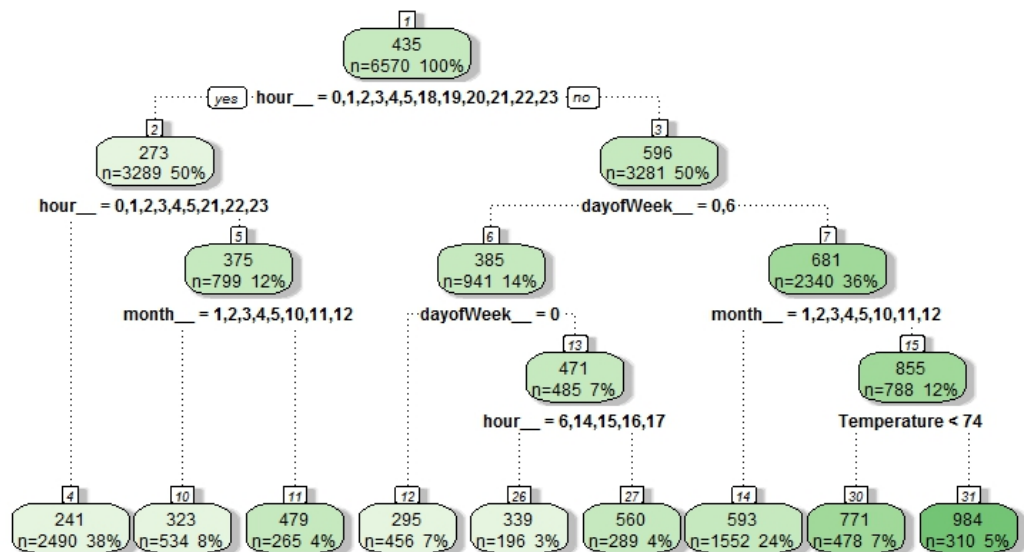
- Prune the tree to the desired size using `prune(fit, cp=)`
- `printcp()` is used to examine the cross-validated error results, select the complexity parameter associated with minimum error, and place it into the `prune()` function

The prune tree results for our model is pasted below:

```

1 install.packages("rpart.plot")
2 install.packages("rpart")
3 library(rpart)
4 library(rattle)
5
6
7 rt <- rpart(power_ ~ Temperature + month_ + day_ + hour_ + dayofWeek_ , data=train)
8
9 test.pred.rtree <- predict(rt,test)
10
11
12 printcp(rt)
13 min.xerror <- rt$cptable[which.min(rt$cptable[, "xerror"]), "CP"]
14 min.xerror
15 rt.pruned <- prune(rt, cp = min.xerror) |
16 fancyRpartPlot(rt.pruned)
17

```



Rattle 2016-Oct-28 18:18:44 310250154

- This tree is giving an overview of decision rules for predicting our outcome.
- We see that hour has the maximum weightage and hence it is the root of the decision/prune tree.
- Subsequent levels explain the weightage in the model respectively.

Plots based on the best fit model:

